

Apache Pig

# What is Pig ?

- A High Level data flow language
- Best for analyzing Large datasets
- Geared for parallel processing
- Compiler produces a sequence of MR programs (in -x mapreduce)
- Operates on files on Local machine/HDFS
- Metadata is not required, but can be used when available (HCatlog)
- Optimal tool for semi-structured data/ETL needs on Hadoop
- Can also be run on Tez & Spark

# Why Pig ?

- Pig follows multi-query approach, which helps combine multiple operation together and hence reducing the number of scans on data
- Its simple and concise: 1/20<sup>th</sup> the lines of code and 1/16<sup>th</sup> the dev time compared to MapReduce
- This Pig eats anything: Structured, Semi-Structured and UnStructured. It adds the complex data-types like Bags, Tuples etc., missing in MR
- Its inherently Lazy! – hence can produce efficient execution plan
- Sampling in any use case. Ex: Data Profiling
- Functionalities can be easily extended through UDFs – including Python, Java, Ruby etc.,

# Pig v/s Hive

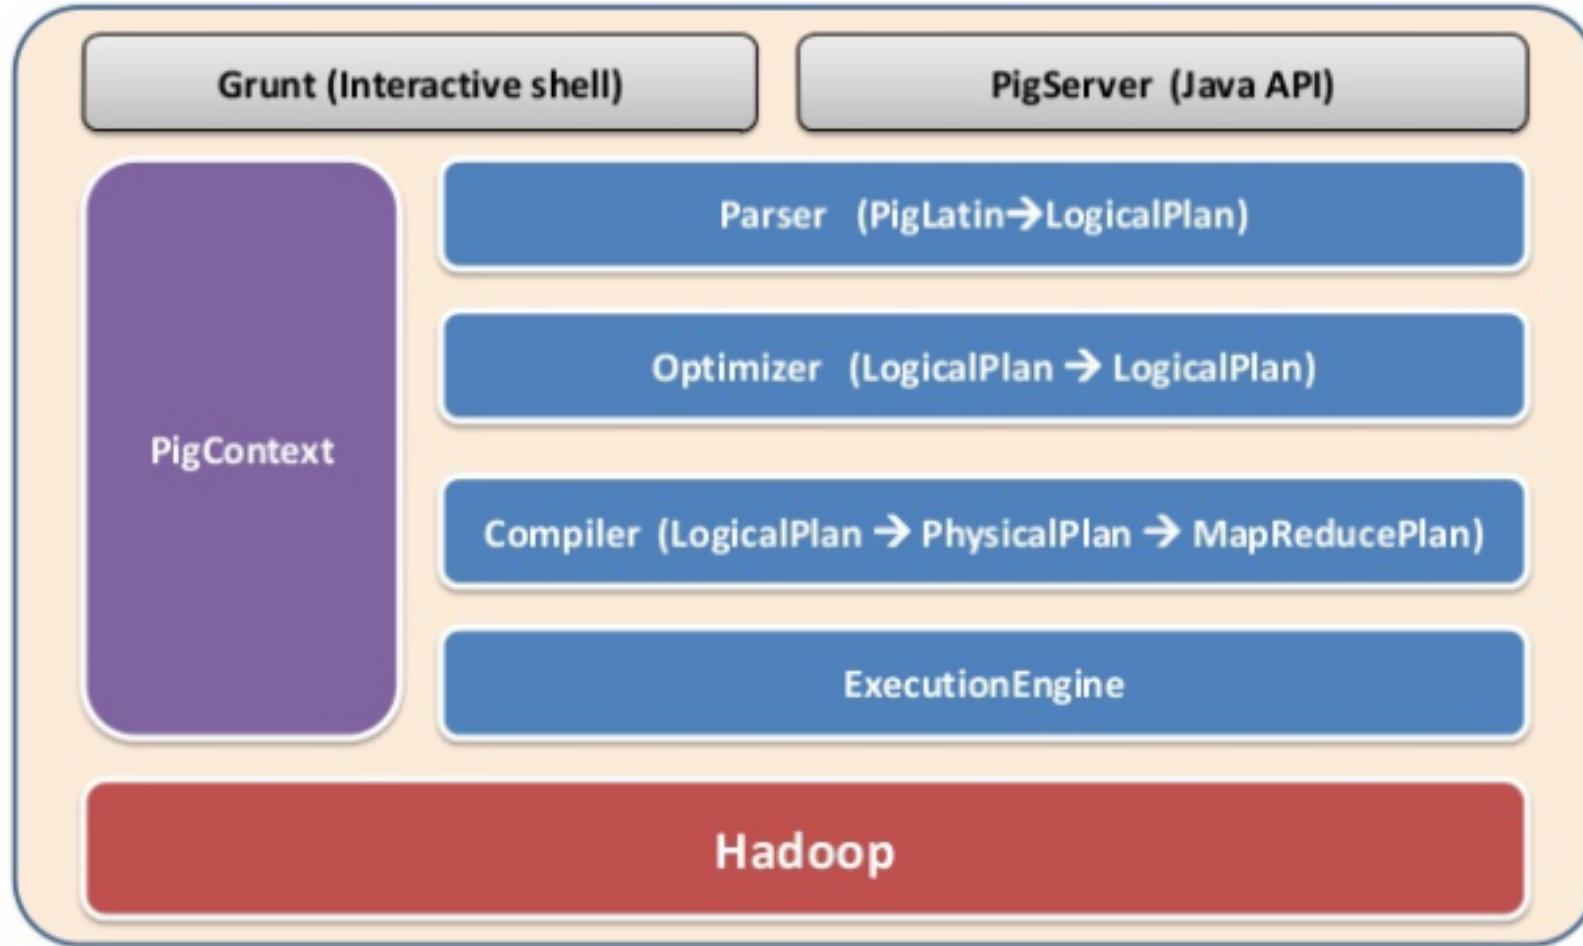
- Hive is SQL like; Pig too – but not really! It's a procedural language
- Pig is used for data analysis and not mainly for creating reports and dashboards like Hive
- They have similar execution times +/- 25% variance
- Pig doesn't need/doesn't have a dedicated metastore unlike Hive. However, it could use one when available.
- Pig operates on client side - Hive on server side
- Pig doesn't support partitions – Hive does
- Pig doesn't support ODBC/JDBC – Hive does

\* If you're good at SQL; chances are you'll not use Pig unless needed to ETL unstructured data

# ‘COMPLEX’ Datatypes

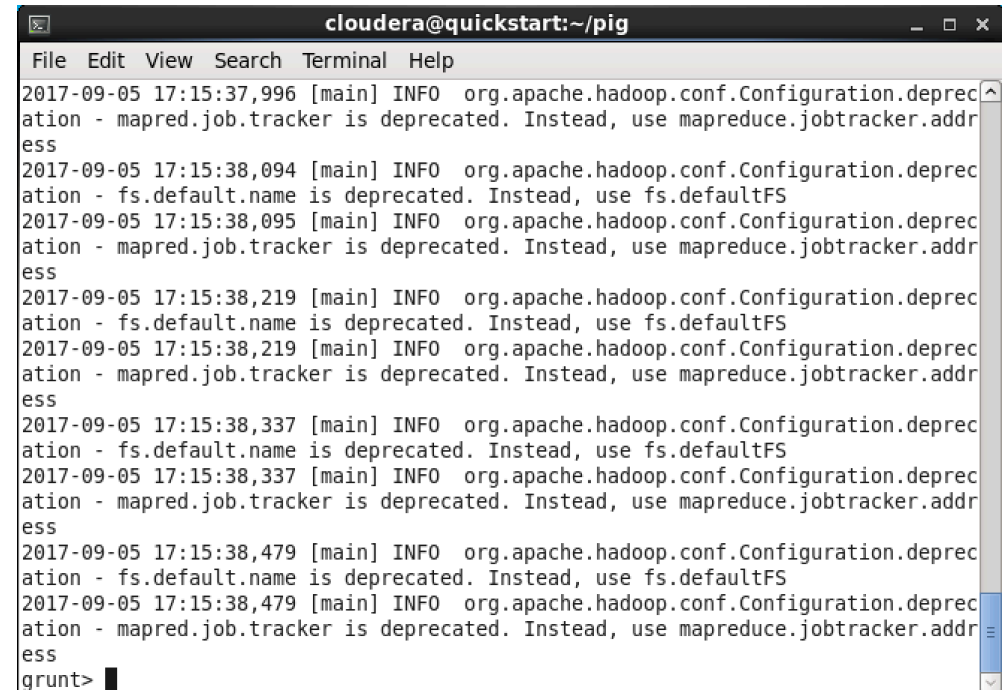
- BAG is collection of TUPLES ~ Tables in our world. Interchangeably called as a RELATION. Ex: ((Amar, 30, M) ,(Vinay, 25, M))
- A TUPLE is an ordered set. Ex: (name, age, sex)
- A FIELD is a piece of data. Ex: `name`
- And then there is MAP. Ex: [‘Name’ # ‘Joydeep’]
- In PIG, Null is UNKNOWN or NON-EXISTENT, similar to SQL though
- There is no string or char; its CHARARRAY

# Architecture of Pig



# Starting PIG

- MapReduce: >pig or >pig -x mapreduce
- Local: >pig -x local
- Tez: >pig -x tez
- Spark: >pig -x spark. 1)doesn't make sense 2)still not completely available
- You should see grunt> shell



```
cloudera@quickstart:~/pig
File Edit View Search Terminal Help
2017-09-05 17:15:37,996 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-05 17:15:38,094 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-05 17:15:38,095 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-05 17:15:38,219 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-05 17:15:38,219 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-05 17:15:38,337 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-05 17:15:38,337 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-09-05 17:15:38,479 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2017-09-05 17:15:38,479 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
grunt>
```

# Loading Data

- LOAD

```
|grunt> a = load '/home/cloudera/pig/flatten.ex' using PigStorage(' ') as (f1:bag{}, f2:chararray, f3:int);
```



# Viewing data (like select)

- DUMP

```
|grunt> dump a;
```

```
|{(1,2)},A,1)  
|(5,6)},B,2)  
|{(7,8),(1,2)},C,3)  
|{(5,2),(1,4)},D,4)  
|{(1,2),(5,7)},A,1)  
|{(1,2),(5,6)},B,2)  
grunt> █
```

# Storing Data

- STORE `|grunt> store c into '/home/cloudera/pig/c';`

```
[cloudera@quickstart pig]$ pwd
/home/cloudera/pig
[cloudera@quickstart pig]$ ls
a.txt  b.txt  c  flatten.ex  pig_1504657295154.log
[cloudera@quickstart pig]$ cd c
[cloudera@quickstart c]$ ls
part-r-00000  _SUCCESS
[cloudera@quickstart c]$ cat part-r-00000
1      2      A      1
1      2      B      2
1      2      C      3
1      4      D      4
5      2      D      4
5      6      B      2
5      7      A      1
7      8      C      3
[cloudera@quickstart c]$
```

# EVERYTHING IN BETWEEN

- Pig is a 'Data Flow' language. Implies, each line of code is a step in the pipeline.
- download script and data files from <https://github.com/ravikodi1/pig>
- `exec /home/cloudera/pig/script.pig` once the extracted folder is placed under `/home/cloudera`

# Important Functions

- FLATTEN
  - TOKENIZE
  - COUNT vs COUNT\_STAR
  - PigStorage() vs PigDump() vs BinStorage()
  - TOBAG
- 
- supports compression for .gz and .bz

# Important Relational Operators

- FOREACH
- SAMPLE
- COGROUP
- SPLIT (Pig doesn't support conditional loops; this is as good as it gets)

# UDFs

custom UDFs can be built using Python, Java, C++, Ruby On Rails.

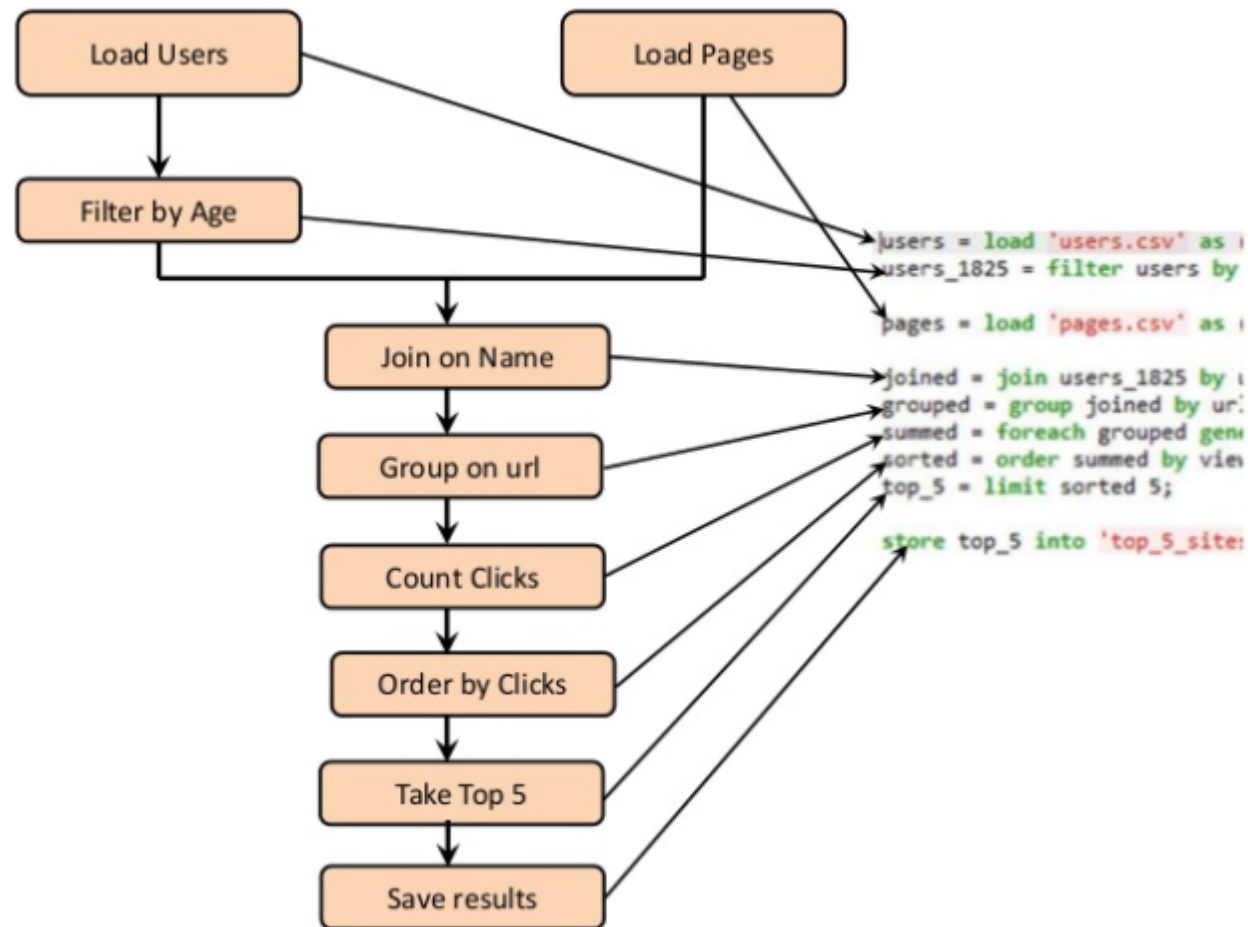
They help extend the functionality of Pig while promoting modularity and reusability of repetitive logic

```
public class TopLevelDomain extends EvalFunc<String> {  
  
    @Override  
    public String exec(Tuple tuple) throws IOException {  
        Object o = tuple.get(0);  
        if (o == null) {  
            return null;  
        }  
        return Validator.getTLD(o.toString());  
    }  
}
```

```

top_5_sites.pig
1 users = load 'users.csv' as (username:chararray, age:int);
2 users_1825 = filter users by age >= 18 and age <= 25;
3
4 pages = load 'pages.csv' as (username:chararray, url:chararray);
5
6 joined = join users_1825 by username, pages by username;
7 grouped = group joined by url;
8 summed = foreach grouped generate group as url, COUNT(joined) as views;
9 sorted = order summed by views desc;
10 top_5 = limit sorted 5;
11
12 store top_5 into 'top_5_sites.csv';
--

```



## Common SQL statements

and their equivalent PIG  
representation

SQL	Pig
...FROM MyTable...	A = LOAD 'MyTable' USING PigStorage('\t') AS (col1:int, col2:int, col3:int);
SELECT col1 + col2, col3 ...	B = FOREACH A GENERATE col1 + col2, col3;
...WHERE col2 > 2	C = FILTER B by col2 > 2;
SELECT col1, col2, sum(col3) FROM X GROUP BY col1, col2	D = GROUP A BY (col1, col2) E = FOREACH D GENERATE FLATTEN(group), SUM(A.col3);
...HAVING sum(col3) > 5	F = FILTER E BY \$2 > 5;
...ORDER BY col1	G = ORDER F BY \$0;
SELECT DISTINCT col1 from X	I = FOREACH A GENERATE col1; J = DISTINCT I;
SELECT col1, count(DISTINCT col2) FROM X GROUP BY col1	K = GROUP A BY col1; L = FOREACH K { M = DISTINCT A.col2; GENERATE FLATTEN(group), count(M); }